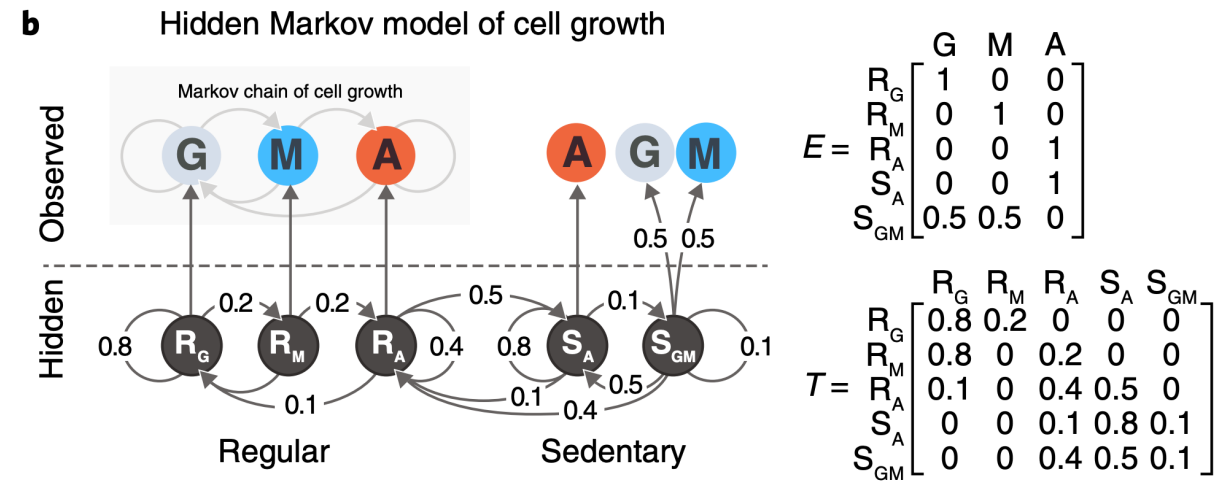
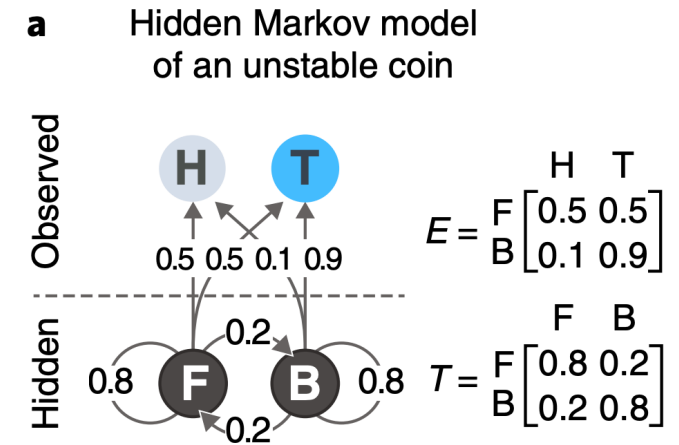


BE 175/275: Machine learning & data-driven modeling in bioengineering

Discussion 7

Topics

- Hidden Markov Model
- Homework: Implementation of HMMs on real-world heart rate data.

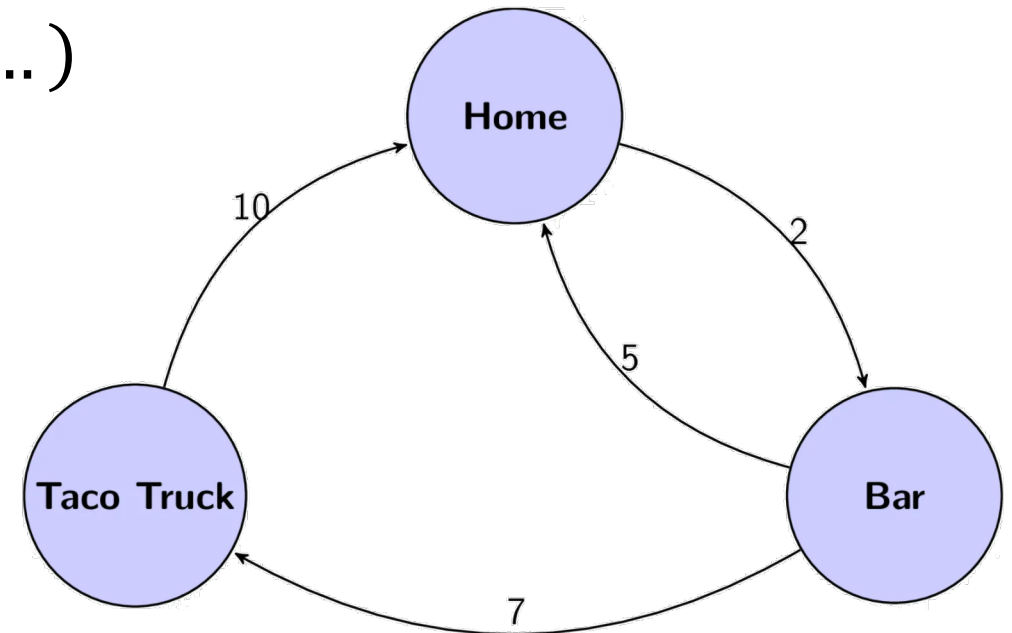


From the [Nature paper](#)

So, what is so “Markov” about it?

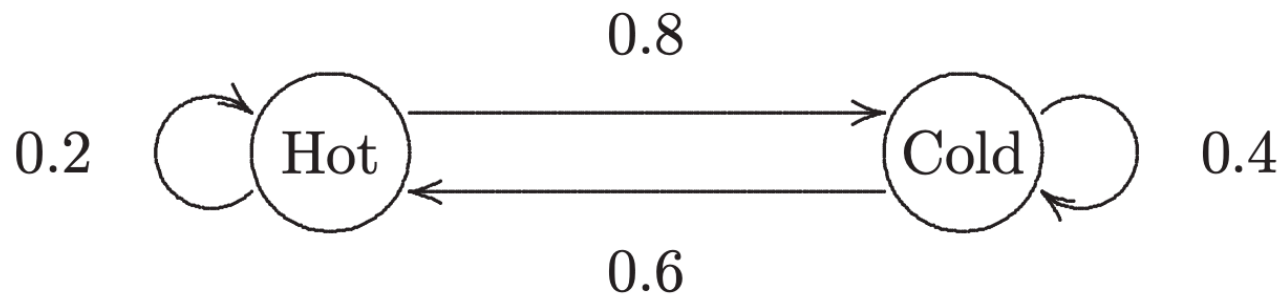
- **Markov Property**: The future depends only on the present state and does not depend on past history.
- $P(X_n | X_{n-1}) = P(X_n | X_{n-1}, X_{n-2}, X_{n-3}, \dots)$

	State Observable	State not fully observable
Autonomous	Markov Chain	Hidden Markov Model
Controlled	Markov Decision Process	Partially observable MDP



Markov Chain

- State at time t is the value of X_t
- State space is the set of values that each X_t can take
- Trajectory is a particular set of values for X_0, X_1, X_2, \dots
- Transition probability $p_{ij} = \mathbb{P}(X_{t+1} = j \mid X_t = i)$.
- Transition matrix $P = (p_{ij})$



$$X_t \begin{cases} \text{Hot} \\ \text{Cold} \end{cases} \begin{matrix} \overbrace{\begin{matrix} & X_{t+1} \\ & \text{Hot} & \text{Cold} \end{matrix}} \\ \left(\begin{array}{cc} 0.2 & 0.8 \\ 0.6 & 0.4 \end{array} \right) \end{matrix}$$

Example from [Rachel Fewster](#)

What is hidden in the Hidden Markov Model?

- HMM as a 5-tuple

- $\lambda(S, O, p, q, \pi)$

- States (S)

- Observations (O)

- Transition probabilities (p)

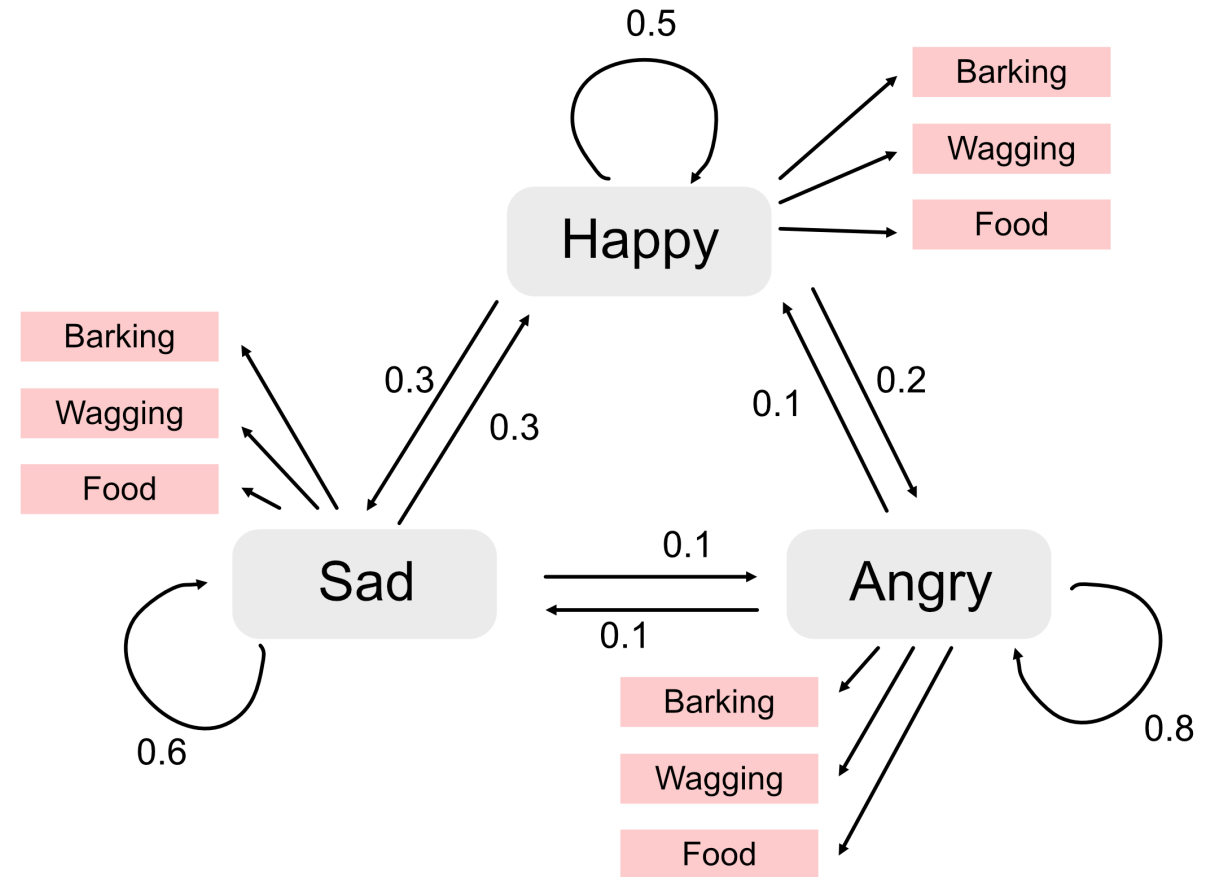
$$p(S_t = j \mid S_{t-1} = i) = p_{ij}$$

- Emission probabilities (q)

$$p(O_t = y \mid S_t = i) = q_i^y$$

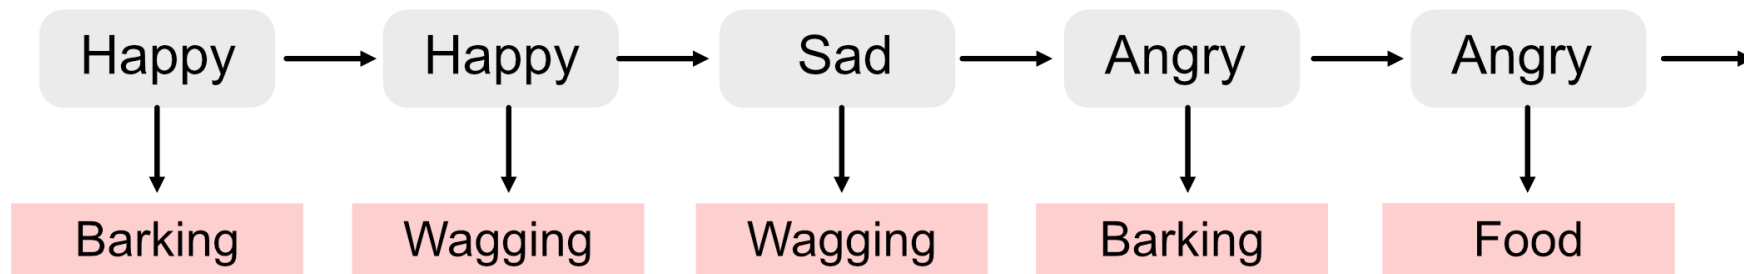
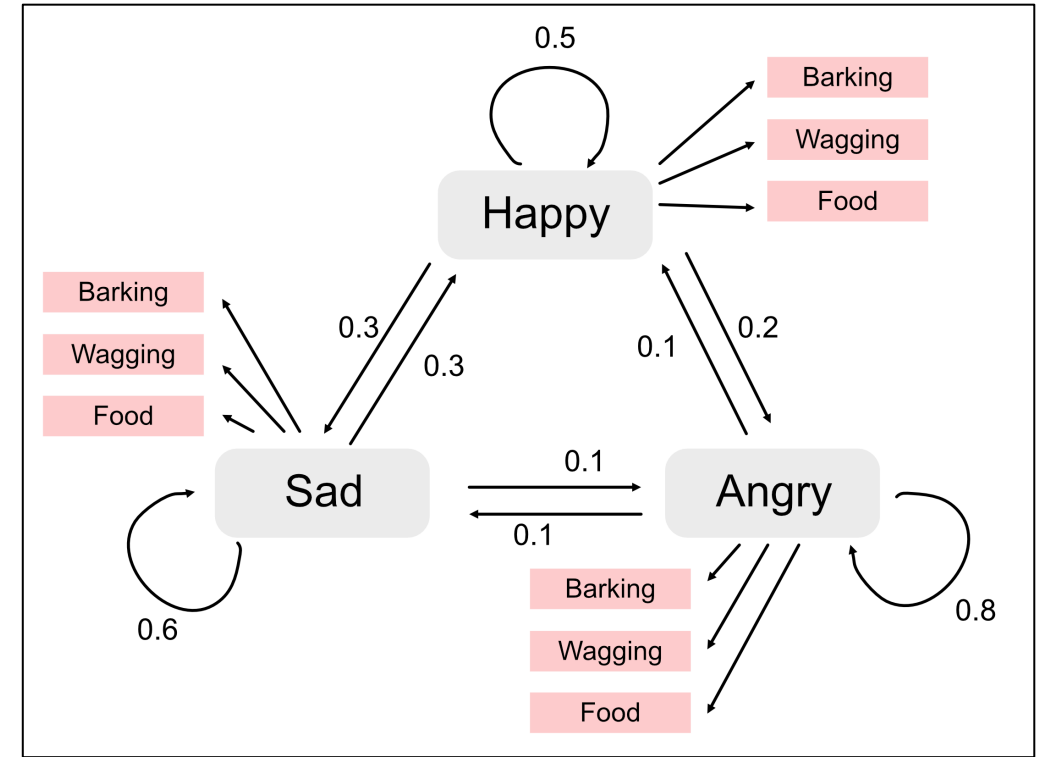
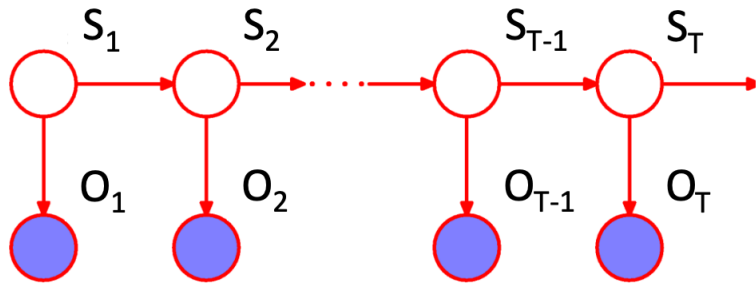
- Initial probabilities (π)

$$p(S_1 = i) = \pi_i$$



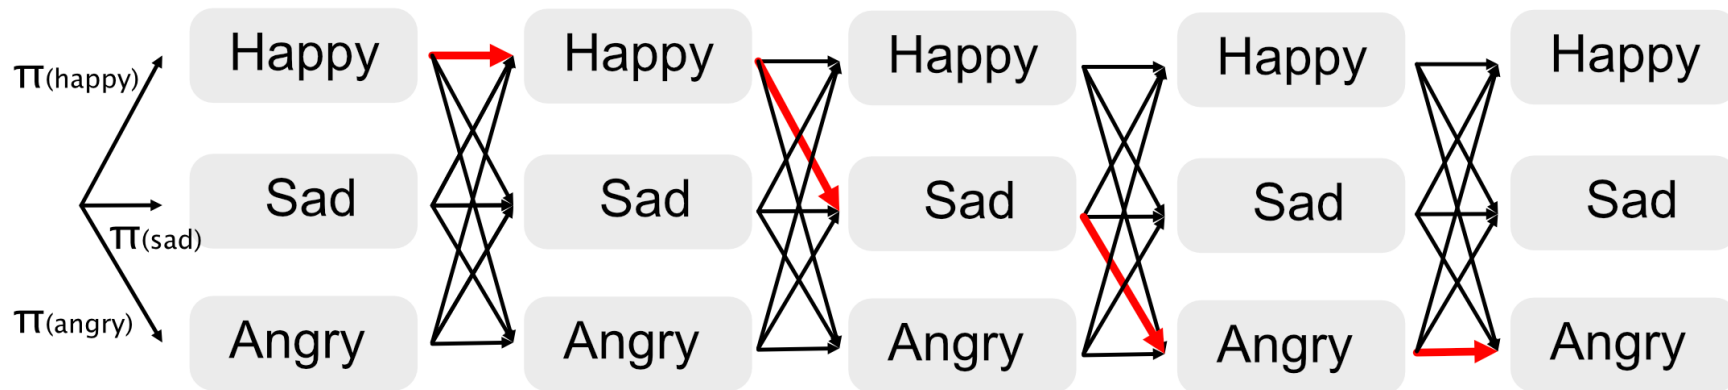
Many diagrams of HMM

- States and observations



- State sequence, observation sequence

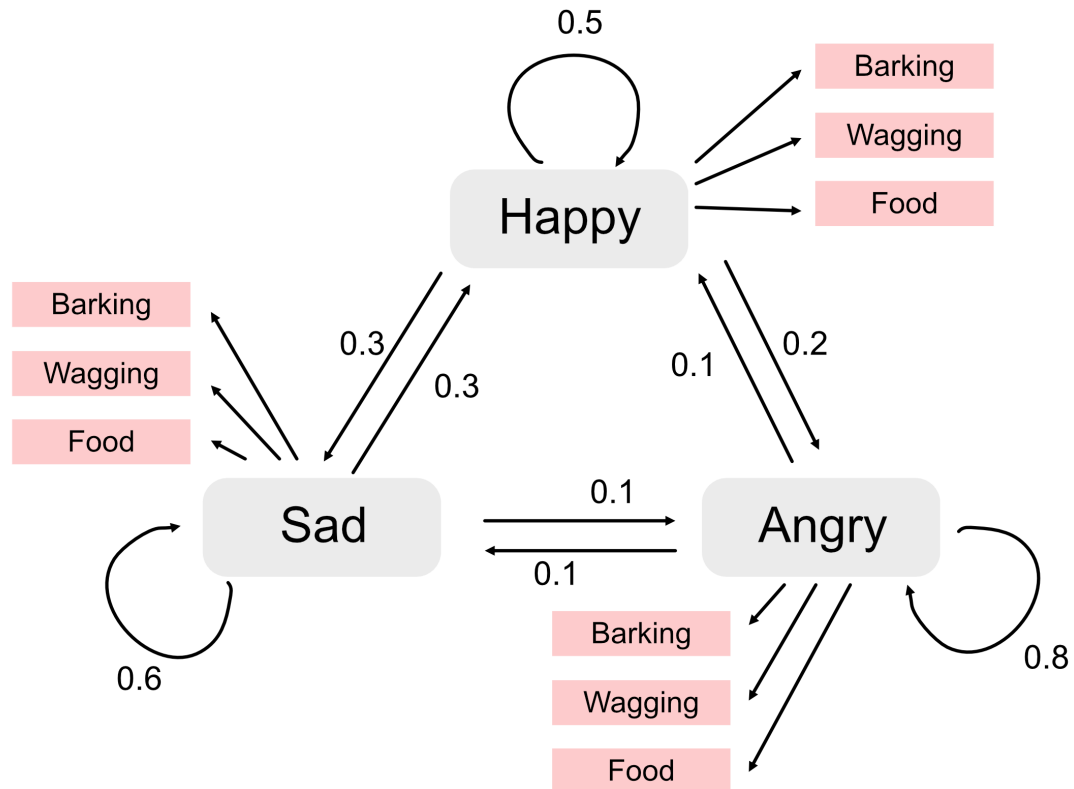
- State transition



Three fundamental problems of HMM

- **Evaluation** Given an HMM $\lambda(p, q)$ and an observation sequence O , determine the likelihood $P(O \mid \lambda)$.
 - **Forward Algorithm**
- **Decoding** Given an observation sequence O and an HMM $\lambda(p, q)$, discover the best hidden state sequence S .
 - One time point: **Forward Backward Algorithm**
 - Sequence: **Viterbi Algorithm**
- **Learning** Given an observation sequence O and the set of states in the HMM, learn the HMM parameters p, q .
 - Expectation Maximization – **Baum-Welch Algorithm**

In class example



- States (S): H, S, A
- Observations (O): B, W, F
- Transition probabilities (p_{ss})

\mathbf{p}	H	S	A	t
H	0.5	0.3	0.1	
S	0.3	0.6	0.1	$t+1$
A	0.2	0.1	0.8	

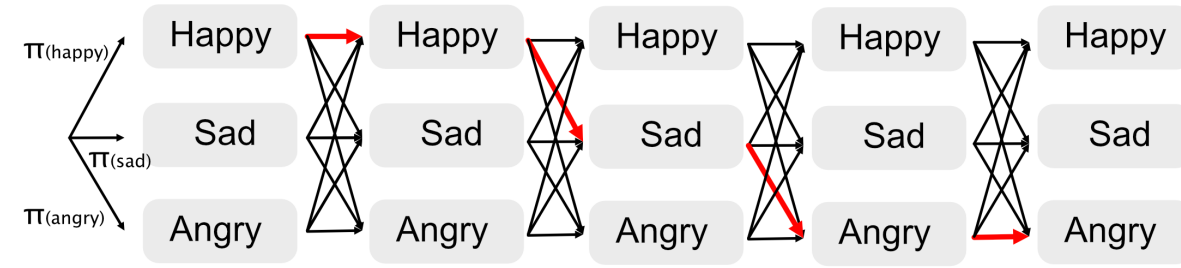
- Emission probabilities (q_s^o)

\mathbf{q}	H	S	A
B	0	0	0.9
W	0.8	1	0
F	0.2	0	0.1

- Initial probability (π_s)

$\boldsymbol{\pi}$	H	S	A
	1	0	0

Dynamic Programming



- Expedite computation by storing the result of previous calculations (Memoization)
- In Forward-Backward algorithm
 - The probability of observing this series before/after current time point AND the current state is k

$$\begin{aligned}
 p(S_t = k, \{O_t\}_{t=1}^T) &= p(O_1, \dots, O_t, S_t = k, O_{t+1}, \dots, O_T) \\
 &= \underbrace{p(O_1, \dots, O_t, S_t = k)}_{\alpha_t^k} \underbrace{p(O_{t+1}, \dots, O_T | S_t = k)}_{\beta_t^k}
 \end{aligned}$$

Compute recursively

- In Viterbi algorithm

Forward Algorithm

- Find the probability of observing “wagging, food, barking” in the first three days?

- Initialize: $\alpha_1^k = p(O_1 | S_1 = k)p(S_1 = k) \quad \forall k$
- Iterate: for $t = 2, \dots, T$

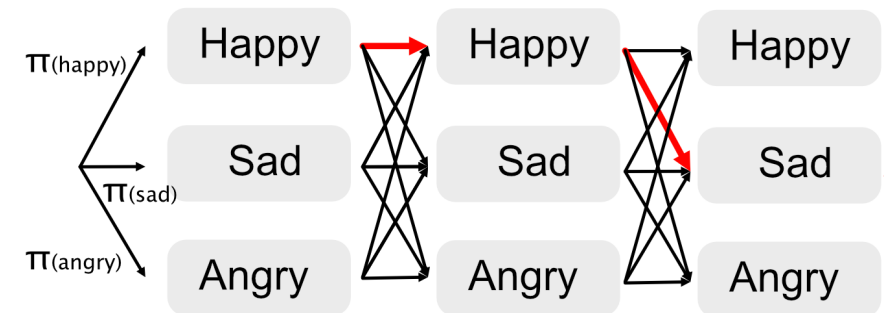
$$\alpha_t^k = p(O_t | S_t = k) \sum_i \alpha_{t-1}^i p(S_t = k | S_{t-1} = i) \quad \forall k$$

- Termination:

$$p(\{O_t\}_{t=1}^T) = \sum_k \alpha_T^k$$

- Saving α along the way

p	H	S	A
H	0.5	0.3	0.1
S	0.3	0.6	0.1
A	0.2	0.1	0.8
q	H	S	A
B	0	0	0.9
W	0.8	1	0
F	0.2	0	0.1
π	H	S	A
	1	0	0



Forward-Backward Algorithm

- If you observe “wagging, food, barking, food” in the first four days, what is the probability that it is **angry** at day 3?

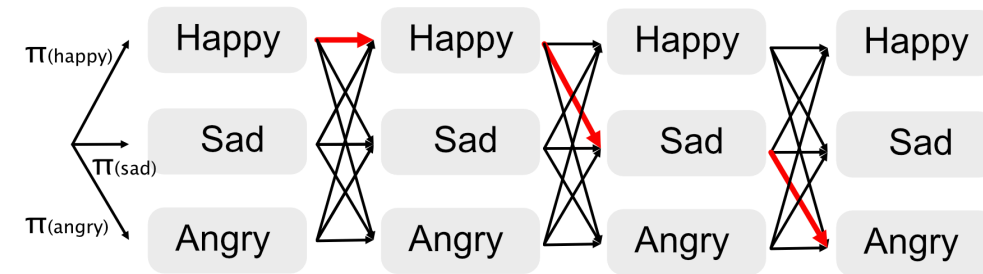
- Initialize: $\beta_T^k = 1 \quad \forall k$
- Iterate: for $t = T-1, \dots, 1$

$$\beta_t^k = \sum_i p(S_{t+1} = i \mid S_t = k) p(O_{t+1} \mid S_{t+1} = i) \beta_{t+1}^i \quad \forall k$$

- Termination: $p(S_t = k, \{O_t\}_{t=1}^T) = \alpha_t^k \beta_t^k$

$$p(S_t = k \mid \{O_t\}_{t=1}^T) = \frac{p(S_t = k, \{O_t\}_{t=1}^T)}{p(\{O_t\}_{t=1}^T)} = \frac{\alpha_t^k \beta_t^k}{\sum_i \alpha_t^i \beta_t^i}$$

p	H	S	A
H	0.5	0.3	0.1
S	0.3	0.6	0.1
A	0.2	0.1	0.8
q	H	S	A
B	0	0	0.9
W	0.8	1	0
F	0.2	0	0.1
π	H	S	A
	1	0	0



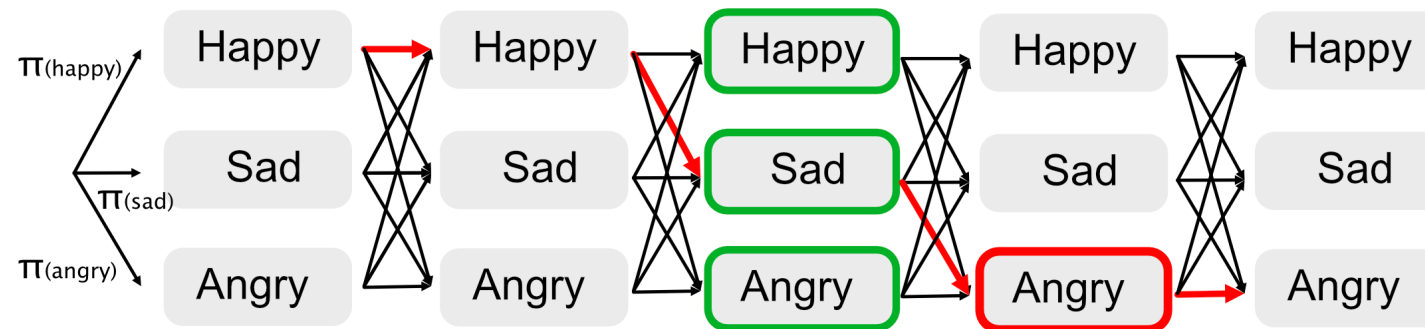
Viterbi algorithm

- If you observe “wagging, food, barking” in the first three days, what are the most likely emotion states of your dog these days?

- Translate to math:

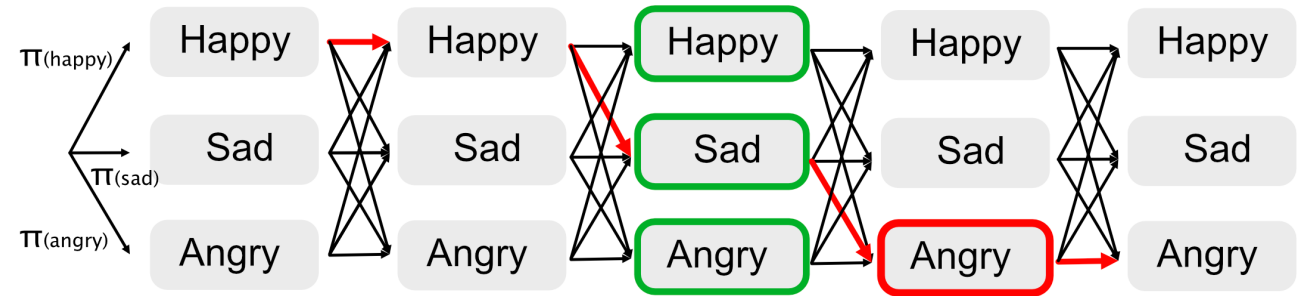
$$\arg \max_{\{S_t\}_{t=1}^T} p(\{S_t\}_{t=1}^T | \{O_t\}_{t=1}^T)$$

- Intuition: How do we find the “shortest” (optimal) path from beginning to the red node?



- Viterbi path

Viterbi algorithm



The optimal path
from source to
state k at t

=

The optimal path
from source to
state i at t-1

×

Transition from
state l to state k

- when we choose the best i
- Define V_t^k as the probability of the most probable Viterbi path of $O_1 O_2 \dots O_t$ ends at state k then we have

$$V_1^k = p(O_1 | S_1 = k) p(S_1 = k)$$

$$V_t^k = p(O_t | S_t = k) \max_i p(S_t = k | S_{t-1} = i) V_{t-1}^i$$

- Eventually, traceback each step to find the path

Baum-Welch Algorithm

- We have a series of observation for our dog, what are the parameters of our HMM?
- Challenge: we have two unknowns

- How does the HMM assign states

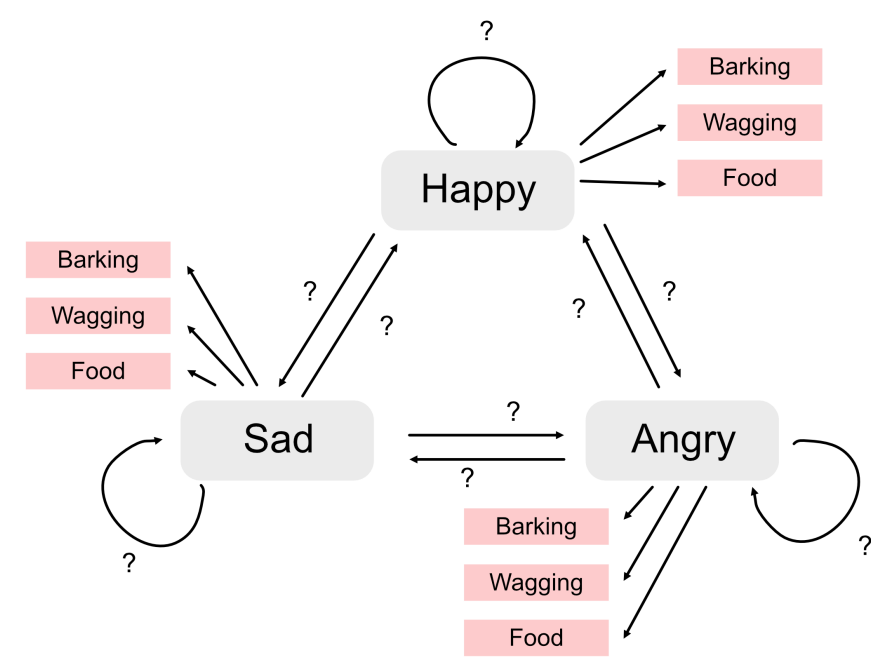
E-step If we know the parameters, we can use forward-backward to assign states given observation

- The parameters in this HMM

M-step If we know the states, we can estimate the parameters by MLE (counting!)

$$p_{ij} = \frac{\sum_{t=1}^{T-1} \xi_{ij}(t)}{\sum_{t=1}^{T-1} \gamma_i(t)} \quad q_i^k = \frac{\sum_{t=1}^T \delta_{O_t=k} \gamma_i(t)}{\sum_{t=1}^T \gamma_i(t)}$$

- Expectation-Maximization (EM) algorithm!
 - Recursive, local optimization



Homework: Implementation of HMMs on real-world heart rate data

- Take heart rate data from one day of an individual and predict which state the person is in over the course of a day.
- Install the HMM package using `pip install hmmlearn`
- Set 3 states: `model = hmmlearn.hmm.GaussianHMM(n_components = 3)`
- (3) How do we know a dataset is normally distributed?
- (6) Fit HMM. Use the transformation you validated in (3)
 - `model.fit(data)`
 - `states = model.predict(data)`
- (7) Use HMM to predict on your moving window averaged data